

Udacity MLND P4: Train a Smart Cab

P.S: I noticed that the newest starting code for this project has been changed for the reward functions. Previously if action is null (i.e. stay as is) the reward is 1, but now it is 0. And before the agent gets 0.5 reward if the move is valid but not what the planner says, now this reward is changed to -0.5. I think these two changes make a lot sense, since the ultimate goal is to let the agent behave what the planner tells it to do, but not violating the traffic rules at the same moment.

Implement a basic driving agent

Implement the basic driving agent, which processes the following inputs at each time step:

- Next waypoint location, relative to its current location and heading,
- Intersection state (traffic light and presence of cars), and,
- Current deadline value (time steps remaining),

And produces some random move/action (`None`, `'forward'`, `'left'`, `'right'`). Don't try to implement the correct strategy! That's exactly what your agent is supposed to learn.

Run this agent within the simulation environment with `enforce_deadline` set to `False` (see `run` function in `agent.py`), and observe how it performs. In this mode, the agent is given unlimited time to reach the destination. The current state, action taken by your agent and reward/penalty earned are shown in the simulator.

In your report, mention what you see in the agent's behavior. Does it eventually make it to the target location?

Answer: The agent will pick a random action from the 4 options, not caring about any other rules or instructions at all. Given unlimited time, the agent will eventually make it to the target location. The problem of this randomized policy is that first the agent ignores all traffic rules, so at times there will be traffic rule violation, even causing accident. Second, it also

ignores the instructions from the planner, which would make it very difficult to get to the target location before the deadline.

Identify and update state

Identify a set of states that you think are appropriate for modeling the driving agent. The main source of state variables are current inputs, but not all of them may be worth representing. Also, you can choose to explicitly define states, or use some combination (vector) of inputs as an implicit state.

At each time step, process the inputs and update the current state. Run it again (and as often as you need) to observe how the reported state changes through the run.

Justify why you picked these set of states, and how they model the agent and its environment.

Answer:

Initially I tried to put all variables from the inputs, i.e. [self.next_waypoint, self.env.sense(), deadline]. But this is a 6-dimensional model, and due to the “Curse of Dimensionality” it would take huge amount of space and time for the training. Then I realized that some of the inputs are not necessary.

The inputs[‘right’] is not necessary, since the agent doesn’t have to worry about who is on its right at all. If the light is red, and the agent’s action is right, it only needs to care about inputs[‘left’]. If the light is red, and the agent’s action is not right, then it is not allowed to move or will get a negative reward. If the light is green, the agent can move as long as its action is not left, otherwise it needs to check what is the situation from inputs[“oncoming”]. In any case, the agent doesn’t have to check inputs[‘right’] at all.

But 5-dimension is still high! Then I noticed that the ultimate goal of the learning agent is as follows: it needs to do whatever planner tells it to do, but not violating the traffic rules, except that when the deadline is coming close, it may violate the rules to follow the planner’s instruction to reach to the destination in time. Therefore, I constructed a new binary variable called “action_okay” that is defined in the following code:

```
def is_action_okay(self, inputs):
```

```
    action_okay = True
    if self.next_waypoint == 'right':
        if inputs['light'] == 'red' and inputs['left'] == 'forward':
            action_okay = False
    elif self.next_waypoint == 'forward':
        if inputs['light'] == 'red':
            action_okay = False
    elif self.next_waypoint == 'left':
        if inputs['light'] == 'red' or (inputs['oncoming'] == 'forward' or inputs['oncoming'] ==
'right'):
            action_okay = False
    return action_okay
```

With this “action_okay” and the way_point and deadline, we have reduced the dimensionality to 3. Further, the original deadline variable is an integer, we need to convert this state to a binary state, such that if the deadline is less than a threshold (say, 3) its value is True, otherwise it is False.

Finally, we have a 3-dimensional state model with the cardinality as follows: 3 for way_point, 2 for action_okay, and 2 for deadline.

Implement Q-Learning

Implement the Q-Learning algorithm by initializing and updating a table/mapping of Q-values at each time step. Now, instead of randomly selecting an action, pick the best action available from the current state based on Q-values, and return that.

Each action generates a corresponding numeric reward or penalty (which may be zero). Your agent should take this into account when updating Q-values. Run it again, and observe the behavior.

What changes do you notice in the agent's behavior?

Answer: The basic Q Learning algorithm is fairly standard. The hyper-parameters are the learning rate, the discount factor for future rewards, and the probability of random actions for each step in order to get rid of the local optima. After implementing this, and tuning these hyper-parameters, the agent will, as expected, try to follow the planner and the traffic rules at the same moment. When deadline is close, it may occasionally violate the traffic rule to get to the destination in time and get the 10 points of the reward.

Enhance the driving agent

Apply the reinforcement learning techniques you have learnt, and tweak the parameters (e.g. learning rate, discount factor, action selection method, etc.), to improve the performance of your agent. Your goal is to get it to a point so that within 100 trials, the agent is able to learn a feasible policy - i.e. reach the destination within the allotted time, with net reward remaining positive.

Report what changes you made to your basic implementation of Q-Learning to achieve the final version of the agent. How well does it perform?

Does your agent get close to finding an optimal policy, i.e. reach the destination in the minimum possible time, and not incur any penalties?

Answer: The basic Q learning works fine, but ideally we want a faster learning at the beginning, and slow it down as we approach the optimal point. In other words, we need a large learning rate initially, then slowly reduce it. Thus I use the idea of “learning rate decay” with this formula: $\text{learning_rate} = \text{learning_rate} * \text{decay_factor}$, for each update step. After reaching a lower bound (say, 0.01) we don’t reduce it anymore.

Similarly, we want to large probability for a random move at the beginning, but we need to reduce it to a small value after a while.

The new Q learning would have two new hyper-parameters, the learning rate decay factor, and the probability decay factor. After tuning these parameters (detailed values can be found in code), the agent can reach to the destination in time even before 50 trials. Due to

the stochasticity, sometimes it need more than 50 trials to train, however it will have the desired behavior within 100 trials for most of the times.

In conclusion, the new Q learning model will let the agent get close to finding an optimal policy in the minimum possible time, less than 50 trials.

Below is the trace of the agent using the learning rate decay factor of 0.8, and the probability decay factor of 0.65. We can see that the agent can learn the correct behavior super fast, and almost always can reach the destination in time even after 10 trials.

Sometimes even if after 50 or more trials, the agent will still miss the deadline occasionally, this is because I noticed that sometime, due to bad luck, the traffic light will delay the agent. But this is very rare.

Simulator.run(): Trial 0

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 1

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 2

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 3

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 4

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 5

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 6

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 7

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 8

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 9

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 10

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 11

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 12

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 13

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 14

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 15

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 16

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 17

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 18

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 19

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 20

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 21

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 22

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 23

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 24

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 25

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 26

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 27

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 28

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 29

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 30

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 31

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 32

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 33

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 34

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 35

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 36

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 37

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 38

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 39

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 40

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 41

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 42

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 43

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 44

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 45

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 46

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 47

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 48

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 49

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 50

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 51

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 52

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 53

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 54

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 55

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 56

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 57

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 58

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 59

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 60

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 61

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 62

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 63

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 64

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 65

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 66

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 67

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 68

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 69

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 70

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 71

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 72

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 73

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 74

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 75

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 76

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 77

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 78

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 79

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 80

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 81

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 82

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 83

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 84

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 85

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 86

Environment.step(): Primary agent ran out of time! Trial aborted.

Simulator.run(): Trial 87

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 88

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 89

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 90

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 91

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 92

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 93

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 94

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 95

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 96

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 97

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 98

Environment.act(): Primary agent has reached destination!

Simulator.run(): Trial 99

Environment.act(): Primary agent has reached destination!